

Clustering

Question 1: We can cluster in one dimension as well as in many dimensions. In this problem, we are going to cluster numbers on the real line. The particular numbers (data points) are 1, 4, 9, 16, 25, 36, 49, 64, 81, and 100, i.e., the squares of 1 through 10. We shall use a k-means algorithm, with two clusters. You can verify easily that no matter which two points we choose as the initial centroids, some prefix of the sequence of squares will go into the cluster of the smaller and the remaining suffix goes into the other cluster. As a result, there are only nine different clusterings that can be achieved, ranging from $\{1\}\{4,9,\dots,100\}$ through $\{1,4,\dots,81\}\{100\}$.

We then go through a reclustering phase, where the centroids of the two clusters are recalculated and all points are reassigned to the nearer of the two new centroids. For each of the nine possible clusterings, calculate how many points are reclassified during the reclustering phase. List out pair of initial centroids that results in *exactly one* point being reclassified.

1. Given data points are 1, 4, 9, 16, 25, 36, 49, 64, 81, 100

There can be 9 different clusters as follows.

$\{1\}, \{4, 9, 16, 25, 36, 49, 64, 81, 100\}$

$\{1, 4\}, \{9, 16, 25, 36, 49, 64, 81, 100\}$

$\{1, 4, 9\}, \{16, 25, 36, 49, 64, 81, 100\}$

$\{1, 4, 9, 16\}, \{25, 36, 49, 64, 81, 100\}$

$\{1, 4, 9, 16, 25\}, \{36, 49, 64, 81, 100\}$

$\{1, 4, 9, 16, 25, 36\}, \{49, 64, 81, 100\}$

$\{1, 4, 9, 16, 25, 36, 49\}, \{64, 81, 100\}$

$\{1, 4, 9, 16, 25, 36, 49, 64\}, \{81, 100\}$

$\{1, 4, 9, 16, 25, 36, 49, 64, 81\}, \{100\}$

2. If we change the centroid only one point has to be shifted b/n the clusters. Let the initial value be 36 and 100. Their mean = 68.

Hence, the clusters will be $\{1, 4, 9, 25, 36, 49, 64\}$ and $\{81, 100\}$.
Centroids of these clusters are 25.5 and 90.5

$$\text{mean} = \frac{25.5 + 90.5}{2} = 58$$

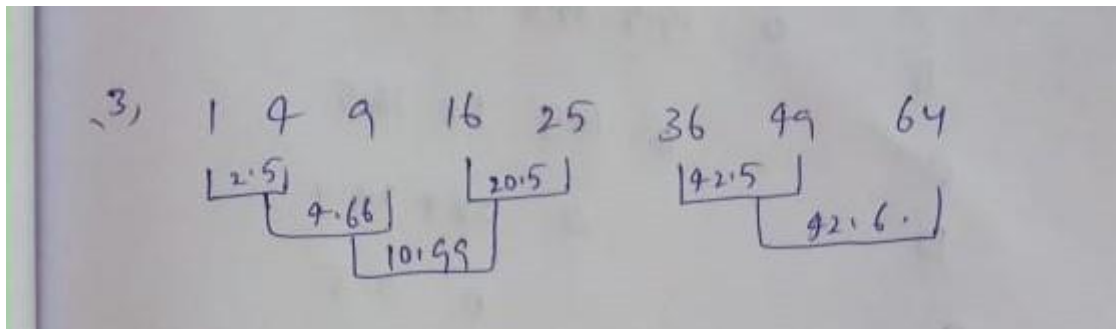
The clusters will be $\{1, 4, 9, 16, 25, 36, 49, 64, 81, 100\}$

Question 2: Suppose we want to assign points to one of two cluster centroids, either $(0,0)$ or $(100,40)$. Depending on whether we use the L_1 or L_2 norm, a point (x,y) could be clustered with a different one of these two centroids. For this problem, you should work out the conditions


2) $(5,5,5)$ is the point which is clustered with Centroid $(0,0)$ when L_1 norm is used and with Centroid $(100,40)$ when L_2 norm is used.

under which a point will be clustered with the centroid $(0,0)$ when the L_1 norm is used, but clustered with the centroid $(100,40)$ when the L_2 norm is used. List out those points.

Question 3: Suppose our data set consists of the perfect squares 1, 4, 9, 16, 25, 36, 49, and 64, which are points in one dimension. Perform a hierarchical clustering on these points, as follows. Initially, each point is in a cluster by itself. At each step, merge the two clusters with the closest centroids, and continue until only two clusters remain. Which centroid of a cluster that exists at some time during this process? Positions are represented to the nearest 0.1



Question 4: Suppose that the true data consists of three clusters, as suggested by the diagram below:



There is a large cluster B centered around the origin (0,0), with 8000 points uniformly distributed in a circle of radius 2. There are two small clusters, A and C, each with 1000 points uniformly distributed in a circle of radius 1. The center of A is at (-10,0) and the center of C is at (10,0).

Suppose we choose three initial centroids x , y , and z , and cluster the points according to which of x , y , or z they are closest. The result will be three *apparent* clusters, which may or may not coincide with the *true* clusters A, B, and C. Say that one of the true clusters is *correct* if there is an apparent cluster that consists of all and only the points in that true cluster. Assuming initial centroids x , y , and z are chosen independently and at random, what is the probability that A is correct? What is the probability that C is correct? What is the probability that both are correct?

4, Given Centroid are x, y, z .

We can assign each of x, y, z to A, B, C in 27 possible ways.

Chance of being in A is $\frac{1000}{10000} = 0.1$

Chance of being in B is $\frac{8000}{10000} = 0.8$

Chance of being in C is $\frac{1000}{10000} = 0.1$


There are six different cases to introduce x, y, z on A, B, C which will be total 27.

The probability that A is correct is 24%.

The probability that C is correct is 24%.

The probability that A & C are correct is 4.8%.

Question 5: Perform a hierarchical clustering of the following six points:



using the *complete-link* proximity measure (the distance between two clusters is the largest distance between any two points, one from each cluster). Find out a cluster at some stage of the agglomeration?

5,

	A	B	C	D	E	F
A						
B	14.1					
C	29.6	46.1				
D	27.4	28.6				
E	15.5	39.5	17.7	18.4		
F						

A and B are low, clustering will be done as follows:

A and B will be clustered with C $\rightarrow AC \rightarrow 29.6$

$BC \rightarrow 15.5$

A and B will be clustered with C $\rightarrow AD \rightarrow 46.6$ and

$BD \rightarrow 32.5$

A and B will be clustered with E $\rightarrow AE \rightarrow 27.4$

$BE \rightarrow 17.7$

A and B will be clustered with F $\rightarrow AF \rightarrow 28.6$

$BF \rightarrow 18.4$

A and D will be clustered $\rightarrow CD - 16.9$

D and F will be clustered $\rightarrow DF - 27.4$

D and E will be clustered $\rightarrow DE \rightarrow 28.6$

CD, DF, DE as CD is low clustering as follows.

C and D will be clustering as C $\rightarrow CE - 17$

and $DE - 28.6$